

8.2. Binary Trees

8.2.1. Binary Trees. A *binary tree* is a rooted tree in which each vertex has at most two children, designated as *left child* and *right child*. If a vertex has one child, that child is designated as either a left child or a right child, but not both. A *full binary tree* is a binary tree in which each vertex has exactly two children or none. The following are a few results about binary trees:

1. If T is a full binary tree with i internal vertices, then T has $i + 1$ terminal vertices and $2i + 1$ total vertices.
2. If a binary tree of height h has t terminal vertices, then $t \leq 2^h$.

More generally we can define a *m-ary tree* as a rooted tree in which every internal vertex has no more than m children. The tree is called a *full m-ary tree* if every internal vertex has exactly m children. An *ordered rooted tree* is a rooted tree where the children of each internal vertex are ordered. A binary tree is just a particular case of *m-ary* ordered tree (with $m = 2$).

8.2.2. Binary Search Trees. Assume S is a set in which elements (which we will call “data”) are ordered; e.g., the elements of S can be numbers in their natural order, or strings of alphabetic characters in lexicographic order. A *binary search tree* associated to S is a binary tree T in which data from S are associate with the vertices of T so that, for each vertex v in T , each data item in the left subtree of v is less than the data item in v , and each data item in the right subtree of v is greater than the data item in v .

Example: Figure 8.2 contains a binary search tree for the set $S = \{1, 2, 3, 4, 5, 6, 7, 8, 9, 10\}$. In order to find a element we start at the root and compare it to the data in the current vertex (initially the root). If the element is greater we continue through the right child, if it is smaller we continue through the left child, if it is equal we have found it. If we reach a terminal vertex without founding the element, then that element is not present in S .

8.2.3. Making a Binary Search Tree. We can store data in a binary search tree by randomly choosing data from S and placing it in the tree in the following way: The first data chosen will be the root of the tree. Then for each subsequent data item, starting at the root we




FIGURE 8.2. Binary Search Tree.

compare it to the data in the current vertex v . If the new data item is greater than the data in the current vertex then we move to the right child, if it is less we move to the left child. If there is no such child then we create one and put the new data in it. For instance, the tree in figure 8.3 has been made from the following list of words choosing them in the order they occur: “IN A PLACE OF LA MANCHA WHOSE NAME I DO NOT WANT TO REMEMBER”.




FIGURE 8.3. Another binary Search Tree.

8.3. Decision Trees, Tree Isomorphisms

8.3.1. Decision Trees. A *decision tree* is a tree in which each vertex represents a question and each descending edge from that vertex represents a possible answer to that question.

Example: The Five-Coins Puzzle. In this puzzle we have five coins C_1, C_2, C_3, C_4, C_5 that are identical in appearance, but one is either heavier or lighter than the others. The problem is to identify the bad coin and determine whether it is lighter or heavier using only a pan balance and comparing the weights of two piles of coins. The problem can be solved in the following way. First we compare the weights of C_1 and C_2 . If C_1 is heavier than C_2 then we know that either C_1 is the bad coin and is heavier, or C_2 is the bad coin and it is lighter. Then by comparing say C_1 with any of the other coins, say C_5 , we can determine whether the bad coin is C_1 and is heavier (if C_1 is heavier than C_5) or it is C_2 and is lighter (if C_1 has the same weight as C_5). If C_1 is lighter than C_2 we proceed as before with “heavier” and “lighter” reversed. If C_1 and C_2 have the same weight we can try comparing C_3 and C_4 in a similar manner. If their weights are the same then we know that the bad coin is C_5 , and we can determine whether it is heavier or lighter by comparing it to say C_1 . The corresponding decision tree is the following:

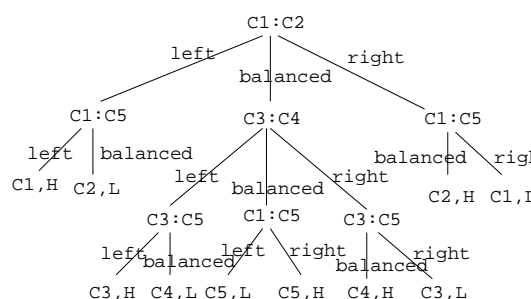


FIGURE 8.4. Decision tree for the 5 coins puzzle.

In each vertex “ $C_i : C_j$ ” means that we compare coins C_i and C_j by placing C_i on the left pan and C_j on the right pan of the balance, and each edge is labeled depending on what side of the balance is heavier. The terminal vertices are labeled with the bad coin and whether it is heavier (H) or lighter (L). The decision tree is optimal in the sense that in the worst case it uses three weighings, and there is no way to solve the problem with less than that—with two weighings we can get

at most nine possible outcomes, which are insufficient to distinguish among ten combinations of 5 possible bad coins and the bad coin being heavier or lighter.

8.3.2. Complexity of Sorting. Sorting algorithms work by comparing elements and rearranging them as needed. For instance we can sort three elements a_1, a_2, a_3 with the decision tree shown in figure 8.5




FIGURE 8.5. Sorting three elements.

Since there are $3! = 6$ possible arrangements of 3 elements, we need a decision tree with at least 6 possible outcomes or terminal vertices. Recall that in a binary tree of height h with t terminal vertices the following inequality holds: $t \leq 2^h$. Hence in our case $6 < 2^h$, which implies $h \geq 3$, so the algorithm represented by the decision tree in figure 8.5 is optimal in the sense that it uses the minimum possible number of comparisons in the worst-case.

More generally in order to sort n elements we need a decision tree with $n!$ outcomes, so its height $h(n)$ will verify $n! \leq 2^{h(n)}$. Since $\log_2(n!) = \Theta(n \log_2 n)$,¹ we have $h(n) = \Omega(n \log_2 n)$. So the worst case complexity of a sorting algorithm is $\Omega(n \log_2 n)$. Since the merge-sort algorithm uses precisely $\Theta(n \log_2 n)$ comparisons, we know that it is optimal.

8.3.3. Isomorphisms of Trees. Assume that T_1 is a tree with vertex set V_1 and T_2 is another tree with vertex set V_2 . If they are rooted trees then we call their roots r_1 and r_2 respectively. We will study three different kinds of tree-isomorphisms between T_1 and T_2 .

¹According to Stirling's formula, $n! \approx n^n e^{-n} \sqrt{2\pi n}$, so taking logarithms $\log_2 n! \approx n \log_2 n - n \log_2 e + \frac{1}{2} \log_2(2\pi n) = \Theta(n \log_2 n)$.

1. Usual graph-isomorphism between trees: T_1 and T_2 are isomorphic if there is a bijection $f : V_1 \rightarrow V_2$ that preserves adjacency, i.e., $f(v)$ is adjacent to $f(w)$ if and only if v is adjacent to w .
2. Rooted-tree-isomorphism: T_1 and T_2 are isomorphic if there is a bijection $f : V_1 \rightarrow V_2$ that preserves adjacency and the root vertex, i.e.:
 - (a) $f(v)$ is adjacent to $f(w)$ if and only if v is adjacent to w .
 - (b) $f(r_1) = r_2$.
3. Binary-tree-isomorphism: Two binary trees T_1 and T_2 are isomorphic if there is a bijection $f : V_1 \rightarrow V_2$ that preserves adjacency, and the root vertex, and left/right children, i.e.:
 - (a) $f(v)$ is adjacent to $f(w)$ if and only if v is adjacent to w .
 - (b) $f(r_1) = r_2$.
 - (c) $f(v)$ is a left child of $f(w)$ if and only if v is a left child of w .
 - (d) $f(v)$ is a right child of $f(w)$ if and only if v is a right child of w .

Example: Figure 8.6 shows three trees which are graph-isomorphic. On the other hand as rooted trees T_2 and T_3 are isomorphic, but they are not isomorphic to T_1 because the root of T_1 has degree 3, while the roots of T_2 and T_3 have degree 2. Finally T_2 and T_3 are not isomorphic as binary trees because the left child of the root in T_2 is a terminal vertex while the left child of the root of T_3 has two children.




FIGURE 8.6. Trees with different kinds of isomorphisms.

Exercise: Find all non-isomorphic 3-vertex free trees, 3-vertex rooted trees and 3-vertex binary trees. *Answer:* Figure 8.7 shows all 5 non-isomorphic 3-vertex binary trees. As rooted trees T_2 – T_5 are isomorphic, but T_1 is not isomorphic to the others, so there are 2 non-isomorphic 3-vertex rooted trees represented for instance by T_1 and T_2 . All of them

are isomorphic as free trees, so there is only 1 non-isomorphic 3-vertex free tree.




FIGURE 8.7. Non-isomorphic binary trees.

8.3.4. Huffman Codes. Usually characters are represented in a computer with fix length bit strings. *Huffman codes* provide an alternative representation with variable length bit strings, so that shorter strings are used for the most frequently used characters. As an example assume that we have an alphabet with four symbols: $A = \{a, b, c, d\}$. Two bits are enough for representing them, for instance $a = 11$, $b = 10$, $c = 01$, $d = 00$ would be one such representation. With this encoding n -character words will have $2n$ bits. However assume that they do not appear with the same frequency, instead some are more frequent than others, say a appears with a frequency of 50%, b 30%, c 15% and d 5%. Then the following encoding would be more efficient than the fix length encoding: $a = 1$, $b = 01$, $c = 001$, $d = 000$. Now in average an n -character word will have $0.5n$ a 's, $0.3n$ b 's, $0.15n$ c 's and $0.05n$ d 's, hence its length will be $0.5n \cdot 1 + 0.3n \cdot 2 + 0.15n \cdot 3 + 0.05n \cdot 3 = 1.7n$, which is shorter than $2n$. In general the length per character of a given encoding with characters a_1, a_2, \dots, a_n whose frequencies are f_1, f_2, \dots, f_n is

$$\frac{1}{F} \sum_{k=1}^n f_k l(a_k),$$

where $l(a_k) = \text{length of } a_k$ and $F = \sum_{k=1}^n f_k$. The problem now is, given an alphabet and the frequencies of its characters, find an optimal encoding that provides minimum average length for words.

Fix length and Huffman codes can be represented by trees like in figure 8.8. The code of each symbol consists of the sequence of labels of the edges in the path from the root to the leaf with the desired symbol.




FIGURE 8.8. Fix length code and Huffman code.

8.3.5. Constructing an Optimal Huffman Code. An optimal Huffman code is a Huffman code in which the average length of the symbols is minimum. In general an optimal Huffman code can be made as follows. First we list the frequencies of all the codes and represent the symbols as vertices (which at the end will be leaves of a tree). Then we replace the two smallest frequencies f_1 and f_2 with their sum $f_1 + f_2$, and join the corresponding two symbols to a common vertex above them by two edges, one labeled 0 and the other one labeled 1. Then common vertex plays the role of a new symbol with a frequency equal to $f_1 + f_2$. Then we repeat the same operation with the resulting shorter list of frequencies until the list is reduced to one element and the graph obtained becomes a tree.

Example: Find the optimal Huffman code for the following table of symbols:

character	frequency
a	2
b	3
c	7
d	8
e	12

Answer: : The successive reductions of the list of frequencies are as follows:

$$\underbrace{2, 3}_5, 7, 8, 12 \rightarrow \underbrace{5, 7}_{12}, 8, 12 \rightarrow 12, 8, 12$$

Here we have a choice, we can choose to add the first 12 and 8, or 8 and the second 12. Let's choose the former:

$$\underbrace{12, 8, 12}_{20} \rightarrow \underbrace{20, 12}_{32} \rightarrow 32$$

The tree obtained is the following:




FIGURE 8.9. Optimal Huffman code 1.

The resulting code is as follows:

character	code
<i>a</i>	1111
<i>b</i>	1110
<i>c</i>	110
<i>d</i>	10
<i>e</i>	0

The other choice yields the following:

$$\underbrace{12, 8, 12}_{20} \rightarrow \underbrace{20, 12}_{32} \rightarrow 32$$




FIGURE 8.10. Optimal Huffman code 2.

character	code
<i>a</i>	111
<i>b</i>	110
<i>c</i>	10
<i>d</i>	01
<i>e</i>	00

8.3.6. Game Trees. Trees are used in the analysis of some games. As an example we study the following game using a tree: Initially there are two piles with 3 coins and 1 coin respectively. Taking turns two players remove any number of coins from one of the piles. The player that removes the last coin loses. The following tree represents all possible sequences of choices. Each node shows the number of coins in each pile, and each edge represents a possible “move” (choice) from one of the players. The first player is represented with a box and the second player is represented with an circle.

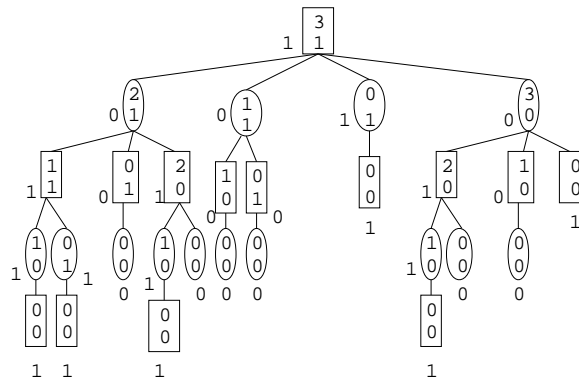


FIGURE 8.11. Tree of a game.

The analysis of the game starts by labeling each terminal vertex with “1” if it represents a victory for the first player and “0” if it represents a victory for the second player. These numbers represent the “value” of each position of the game, so that the first player is interested in making it “maximum” and the second player wants to make it “minimum”. Then we continue labeling the rest of the vertices in the following way. After all the children of a given vertex have been labeled, we label the vertex depending on whether it is a “first player” position (box) or a “second player” position (circle). First player positions are labeled with the maximum value of the labels of its children, second player positions are labeled with the minimum

value of the labels of its children. This process is called the *minimax procedure*. Every vertex labeled “1” will represent a position in which the first player has advantage and can win if he/she works without making mistakes; on the other hand, vertices labeled “0” represent positions for which the second player has advantage. Now the strategy is for the first player to select at each position a children with maximum value, while the second player will be interested in selecting children with minimum value. If the starting position has been labeled “1” that means that the first player has a winning strategy, otherwise the second player has advantage. For instance in the present game the first player has advantage at the initial position, and only one favorable movement at that point: $\binom{3}{1} \rightarrow \binom{0}{1}$, i.e., he/she must remove all 3 coins from the first pile. If for any reason the first player makes a mistake and removes say one coin from the first pile, going to position $\binom{2}{1}$, then the second player has one favorable move to vertex $\binom{0}{1}$, which is the one with minimum “value”.

Alpha-beta pruning. In some games the game tree is so complicated that it cannot be fully analyzed, so it is built up to a given depth only. The vertices reached at that depth are not terminal, but they can be “evaluated” using heuristic methods (for instance in chess usually losing a knight is a better choice than losing the queen, so a position with one queen and no knights will have a higher value than one with no queen and one knight). Even so the evaluation and labeling of the vertices can be time consuming, but we can bypass the evaluation of many vertices using the technique of *alpha-beta pruning*. The idea is to skip a vertex as soon as it becomes obvious that its value will not affect the value of its parent. In order to do that with a first player (boxed) vertex v , we assign it an *alpha value* equal to the maximum value of its children evaluated so far. Assume that we are evaluating one of its children w , which will be a second player (circled) position. If at any point a children of w gets a value less than or equal to the alpha value of v then it will become obvious that the value of w is going to be less than the current alpha value of v , so it will not affect the value of v and we can stop the process of evaluation of w (prune the subtree at w). That is called an *alpha cutoff*. Similarly, at a second player (circled) vertex v , we assign a *beta value* equal to the minimum value of its children evaluated so far, and practice a *beta cutoff* when one of its grandchildren gets a value greater than or equal to the current beta value of v , i.e., we prune the subtree at w , where w is the parent of that grandchildren.




FIGURE 8.12. Alpha cutoff.